

Introduction to CUDA

... curtesy of Dr. Massimo Bernaschi (CNR - <http://www.iac.cnr.it/~massimo>)

Ivan Girotto – igirotto@ictp.it

Information & Communication Technology Section (ICTS)

International Centre for Theoretical Physics (ICTP)



Scuola Internazionale Superiore
di Studi Avanzati



What is CUDA?


- CUDA = **Compute Unified Device Architecture**
 - Expose general-purpose GPU computing as first-class capability
 - Retain traditional DirectX/OpenGL graphics performance
- CUDA C
 - Based on industry-standard C
 - A handful of language extensions to allow heterogeneous programs
 - Straightforward APIs to manage devices, memory, etc.

CUDA Programming Model

- The GPU is viewed as a compute device that:
 - has its own RAM (device memory)
 - runs data-parallel portions of an application as kernels by using many threads
- GPU vs. CPU threads
 - GPU threads are extremely lightweight
 - Very little creation overhead
 - GPU needs 1000s of threads for full efficiency
 - A multi-core CPU needs only a few (basically one thread per core)

CUDA C Jargon: The Basics

- The CPU and its memory (host memory)
- The GPU and its memory (device memory)

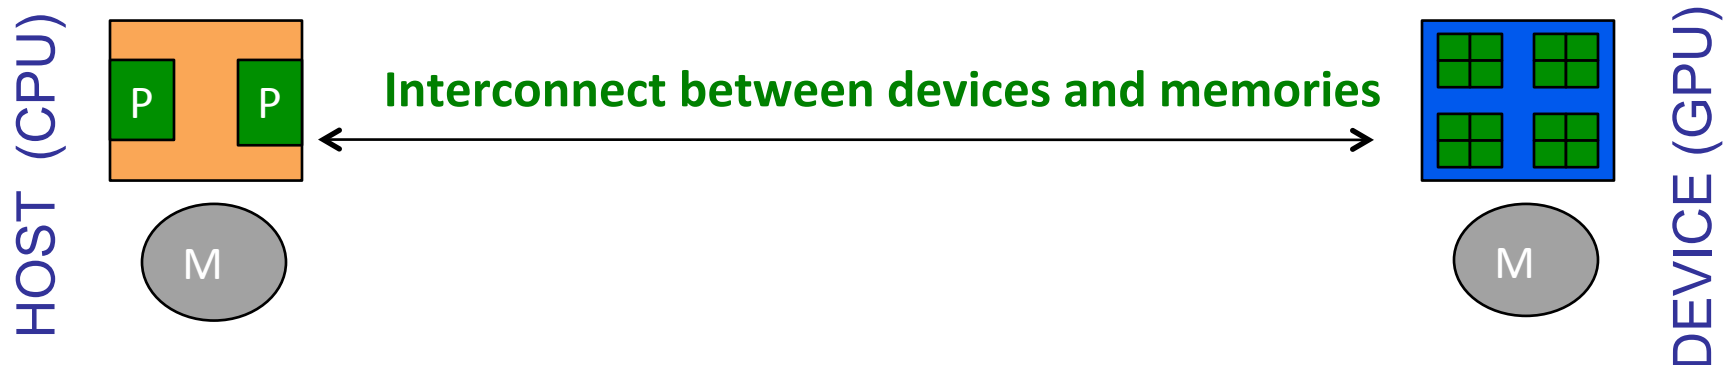


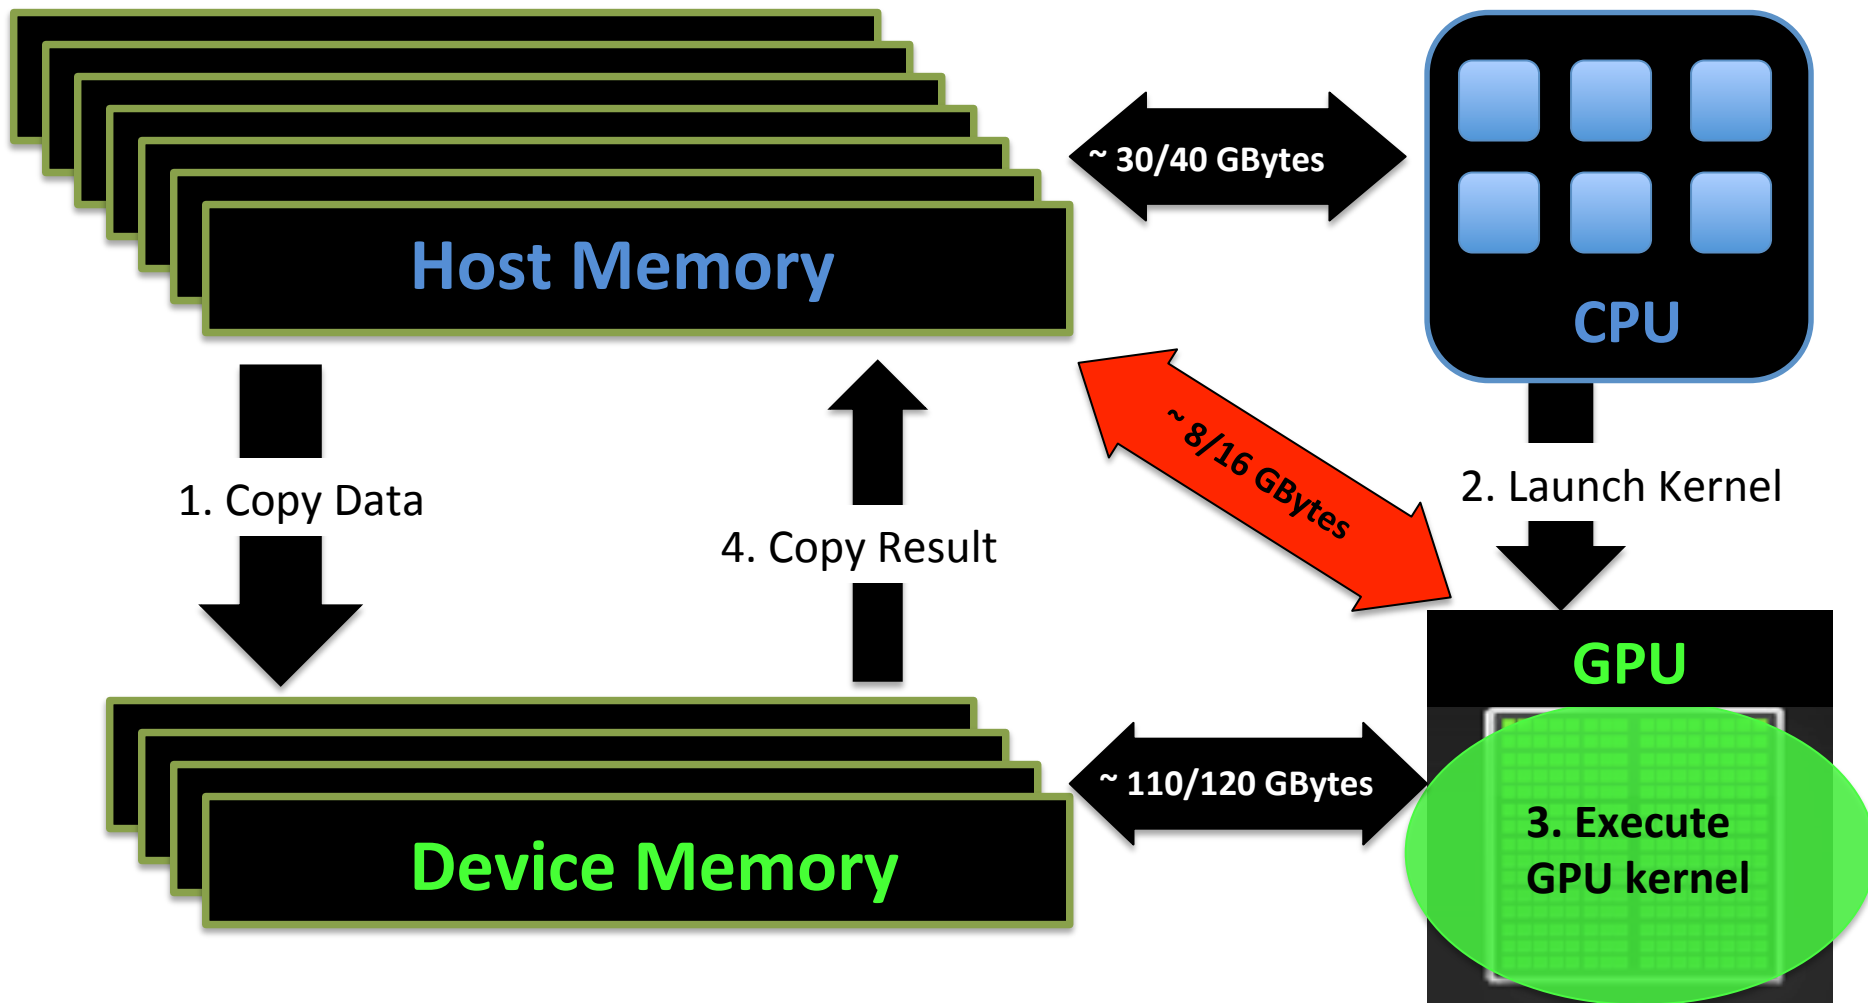
Host



Device

What Programmer Expresses in CUDA





What Programmer Expresses in CUDA

- ✓ Computation partitioning (where does computation occur?)
 - ✓ Declarations on functions `__host__`, `__global__`, `__device__`
 - ✓ Mapping of thread programs to device: **compute <<<gs, bs>>>(<args>)**
- ✓ Data partitioning (where does data reside, who may access it and how?)
 - ✓ Declarations on data `__shared__`, `__device__`, `__constant__`, ...
- ✓ Data management and orchestration
 - ✓ Copying to/from host:
e.g., `cudaMemcpy(h_obj, d_obj, size, cudaMemcpyDeviceToHost)`
- ✓ Concurrency management
 - ✓ e.g. `__syncthreads()`



Hello, World!

```
int main( void ) {  
    printf( "Hello, World!\n" );  
    return 0;  
}
```

- To compile: **nvcc -o hello_world hello_world.cu**
- To execute: **./hello_world**
- This basic program is just standard C that runs on the *host*
- NVIDIA's compiler (**nvcc**) will not complain about CUDA programs with no *device* code
- At its simplest, CUDA C is just C!

Hello, World! with Device Code

```
__global__ void kernel( void ) {  
}  
  
int main( void ) {  
  
    kernel<<<1,1>>>() ;  
    printf( "Hello, World!\n" );  
    return 0;  
}
```

To compile: **nvcc -o simple_kernel simple_kernel.cu**

To execute: **./simple_kernel**

Hello, World! with Device Code

```
__global__ void kernel( void ) {  
}
```

- CUDA C keyword `__global__` indicates that a function
 - Runs on the device
 - Called from host code
- **nvcc** splits source file into host and device components
 - NVIDIA's compiler handles device functions like `kernel()`
 - Standard host compiler handles host functions like `main()`
 - `gcc`, `icc`, ...
 - **Microsoft Visual C**

Hello, World! with Device Code

```
int main( void ) {  
    kernel<<< 1, 1 >>>();  
    printf( "Hello, World!\n" );  
    return 0;  
}
```

- Triple angle brackets mark a call from *host* code to *device* code
 - A “kernel launch” in CUDA jargon
 - We’ll discuss the parameters inside the angle brackets later
- This is all that’s required to execute a function on the GPU!

A More Complex Example

- A kernel to add two integers:

```
__global__ void add( int *a, int *b, int *c ) {  
    *c = *a + *b;  
}
```

- As before, `__global__` is a CUDA C keyword meaning
 - `add()` will execute on the device
 - `add()` will be called from the host

A More Complex Example

- Notice that now we use *pointers* for all our variables:


```
__global__ void add( int *a, int *b, int *c ) {  
    *c = *a + *b;  
}
```

- `add()` runs on the device...so `a`, `b`, and `c` must point to device memory
- How do we allocate memory on the GPU?



Memory Management

- Up to CUDA 4.0 host and device memory were distinct entities from the programmers' viewpoint
 - Device pointers point to GPU memory
 - May be passed to and from host code
 - (In general) May not be dereferenced from host code
 - Host pointers point to CPU memory
 - May be passed to and from device code
 - (In general) May not be dereferenced from device code



Starting on CUDA 4.0 there is a **Unified Virtual Addressing** feature.

Memory Management

- Basic CUDA API for dealing with device memory
 - `cudaMalloc(&p, size), cudaFree(p),
cudaMemcpy(t, s, size, direction)`
 - Similar to their C equivalents: `malloc()`, `free()`,
`memcpy()`

pointer to pointer

```
int main( void ) {
    int a, b, c;                // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
    int size = sizeof( int );   // we need space for an integer
    // allocate device copies of a, b, c
    cudaMalloc( (void**)&dev_a, size );
    cudaMalloc( (void**)&dev_b, size );
    cudaMalloc( (void**)&dev_c, size );
    a = 2;
    b = 7;
    // copy inputs to device
    cudaMemcpy( dev_a, &a, size, cudaMemcpyHostToDevice );
    cudaMemcpy( dev_b, &b, size, cudaMemcpyHostToDevice );
    // launch add() kernel on GPU, passing parameters
    add<<< 1, 1 >>>( dev_a, dev_b, dev_c );
    // copy device result back to host copy of c
    cudaMemcpy( &c, dev_c, size, cudaMemcpyDeviceToHost );
    cudaFree( dev_a ); cudaFree( dev_b ); cudaFree( dev_c )
    return 0;
}
```



```
#include "cuPrintf.cu"
```

```
__global__ void testKernel(int param){  
    cuPrintf("Param value: %d\n", param);  
}
```

```
int main(void){  
    // initialize cuPrintf  
    cudaPrintfInit();  
    int a = 456;  
    testKernel<<<4,1>>>(a);  
    // display the device's greeting  
    cudaPrintfDisplay();  
    // clean up after cuPrintf  
    cudaPrintfEnd();  
} // compile with nvcc -o test.x test.cu -I$CUDADIR/samples/0_Simple/simplePrintf
```



**Also simple
printf()
works!!!**

CUDA Error Checking

- CUDA host function calls usually return a value of type **cudaError_t**

```
cudaError_t cudaMalloc (void **devPtr, size_t size)
```

- Example: to check if device allocation was successful

```
cudaError_t error;  
[...]  
error = cudaMalloc(&d_a, memSize);  
if (error != cudaSuccess)  
{  
    printf("Error in device allocation: %s\n", ! ! ! cudaGetErrorString(error));  
}
```

CUDA Error Checking

- Kernels can't have a return value, so `cudaGetLastError()` is used

```
cudaError_t error;  
[...]  
myKernel<<<1, 1>>>(a_d);  
error = cudaGetLastError();  
if (error != cudaSuccess)  
{  
    printf("Error in Kernel execution: %s\n", cudaGetErrorString(error) );  
}
```

Parallel Programming in CUDA C

- But wait...GPU computing is about **massive** parallelism
- So how do we run code *in parallel* on the device?
- Solution lies in the parameters between the triple angle brackets:

```
add<<< 1, 1 >>>( dev_a, dev_b, dev_c );  
      ↓  
add<<< N, 1 >>>( dev_a, dev_b, dev_c );
```

- Instead of executing **add()** once, **add()** executed **N** times in parallel

Parallel Programming in CUDA C

- With `add()` running in parallel...let's do *vector* addition
- Terminology: Each parallel invocation of `add()` referred to as a *block*
- Kernel can refer to its block's index with the variable `blockIdx.x`
- Each block adds a value from `a[]` and `b[]`, storing the result in `c[]`:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- By using `blockIdx.x` to index arrays, each block handles a different index
- `blockIdx.x` is the first example of a CUDA predefined variable.

Parallel Programming in CUDA C

- We write this code:

```
__global__ void add( int *a, int *b, int *c ) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

- This is what runs in parallel on the device:

Block 0

```
c[0]=a[0]+b[0];
```

Block 1

```
c[1]=a[1]+b[1];
```

Block 2

```
c[2]=a[2]+b[2];
```

Block 3

```
c[3]=a[3]+b[3];
```

Parallel Addition: `main()`

```
#define N 512
int main( void ) {
    int *a, *b, *c;           // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
    int size = N * sizeof( int ); // we need space for 512
                                   // integers

    // allocate device copies of a, b, c
    cudaMalloc( (void**)&dev_a, size );
    cudaMalloc( (void**)&dev_b, size );
    cudaMalloc( (void**)&dev_c, size );

    a = (int*)malloc( size );
    b = (int*)malloc( size );
    c = (int*)malloc( size );

    random_ints( a, N );
    random_ints( b, N );
```

Parallel Addition: `main()` (cont)

```
// copy inputs to device
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );

// launch add() kernel with N parallel blocks
add<<< N, 1 >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( c, dev_c, size, cudaMemcpyDeviceToHost );

free( a ); free( b ); free( c );
cudaFree( dev_a );
cudaFree( dev_b );
cudaFree( dev_c );
return 0;
}
```


Review

- Difference between “host” and “device”
 - Host = CPU
 - Device = GPU
- Using `__global__` to declare a function as device code
 - Runs on device
 - Called from host
- Passing parameters from host code to a device function

Review (cont)

- Basic device memory management
 - `cudaMalloc()`
 - `cudaMemcpy()`
 - `cudaFree()`
- Launching parallel kernels
 - Launch **N** copies of `add()` with: `add <<< N, 1 >>>();`
 - `blockIdx.x` allows to access block's index

Exercise: look at, compile and run the [add_simple_blocks.cu](#) code

Threads

- Terminology: A block can be split into parallel *threads*
- Let's change vector addition to use parallel threads instead of parallel blocks:

```
__global__ void add( int *a, int *b, int *c ) {  
  c[ threadIdx.x ] = a[ blockIdx.x ] + b[ threadIdx.x ];  
}
```

- We use `threadIdx.x` instead of `blockIdx.x` in `add()`
- `main()` will require one change as well...

Parallel Addition (Threads): `main()`

```
#define N 512
int main( void ) {
    int *a, *b, *c;           // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
    int size = N * sizeof( int ); // we need space for 512
                                   // integers

    // allocate device copies of a, b, c
    cudaMalloc( (void**)&dev_a, size );
    cudaMalloc( (void**)&dev_b, size );
    cudaMalloc( (void**)&dev_c, size );

    a = (int*)malloc( size );
    b = (int*)malloc( size );
    c = (int*)malloc( size );

    random_ints( a, N );
    random_ints( b, N );
```

Parallel Addition (Threads): `main()` (cont)

```
// copy inputs to device
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );

// launch add() kernel with N parallel threads
add<<< 1, N >>>( dev_a, dev_b, dev_c );

// copy device result back to host copy of c
cudaMemcpy( c, dev_c, size, cudaMemcpyDeviceToHost );


free( a ); free( b ); free( c );
cudaFree( dev_a );
cudaFree( dev_b );
cudaFree( dev_c );
return 0;
}
```

Using Threads And Blocks

- We've seen parallel vector addition using
 - Many blocks with 1 thread apiece
 - 1 block with many threads
- Let's adapt vector addition to use lots of *both* blocks and threads
- After using threads and blocks together, we'll talk about *why* threads
- First let's discuss data indexing...

Indexing Arrays With Threads & Blocks

- No longer as simple as just using `threadIdx.x` or `blockIdx.x` as indices
- To index array with 1 thread per entry (using 8 threads/block)



- If we have **M** threads/block, a unique array index for each entry is given by


```
int index = threadIdx.x + blockIdx.x * M;
```

```
int index =      x      +      y      * width;
```

Arrows point from `threadIdx.x` to `x`, from `blockIdx.x` to `y`, and from `M` to `width`.

Indexing Arrays: Example

- In this example, the **red** entry would have an index of 21:



Indexing Arrays: other examples (4 blocks with 4 threads *per* block)

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = 7;  
}
```

Output: 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = blockIdx.x;  
}
```

Output: 0 0 0 0 1 1 1 1 2 2 2 2 3 3 3 3

```
__global__ void kernel( int *a )  
{  
    int idx = blockIdx.x*blockDim.x + threadIdx.x;  
    a[idx] = threadIdx.x;  
}
```

Output: 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3

Addition with Threads and Blocks

- **blockDim.x** is a built-in variable for threads per block:

```
int index = threadIdx.x + blockIdx.x * blockDim.x;
```

- **gridDim.x** is a built-in variable for blocks in a grid;
- A combined version of our vector addition kernel to use blocks *and* threads:

```
__global__ void add( int *a, int *b, int *c ) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    c[ index ] = a[ index ] + b[ index ];  
}
```

- So what changes in **main()** when we use both blocks and threads?

Parallel Addition (Blocs/Threads): `main()`

```
#define N    (2048 * 2048)
#define THREADS_PER_BLOCK 512
int main( void ) {
    int *a, *b, *c;           // host copies of a, b, c
    int *dev_a, *dev_b, *dev_c; // device copies of a, b, c
    int size = N * sizeof( int ); // we need space for N integers

    // allocate device copies of a, b, c
    cudaMalloc( (void**)&dev_a, size );
    cudaMalloc( (void**)&dev_b, size );
    cudaMalloc( (void**)&dev_c, size );

    a = (int*)malloc( size );
    b = (int*)malloc( size );
    c = (int*)malloc( size );

    random_ints( a, N );
    random_ints( b, N );
```

Parallel Addition (Threads): `main()` (cont)

```
// copy inputs to device
```

```
cudaMemcpy( dev_a, a, size, cudaMemcpyHostToDevice );
```

```
cudaMemcpy( dev_b, b, size, cudaMemcpyHostToDevice );
```

```
// launch add() kernel with blocks and threads
```

```
add<<< N/THREADS_PER_BLOCK, THREADS_PER_BLOCK >>>(dev_a, dev_b, dev_c);
```

```
// copy device result back to host copy of c
```

```
cudaMemcpy( c, dev_c, size, cudaMemcpyDeviceToHost );
```

```
free( a ); free( b ); free( c );
```

```
cudaFree( dev_a );
```

```
cudaFree( dev_b );
```

```
cudaFree( dev_c );
```

```
return 0;
```

```
}
```

Exercises

- Array reversal: fill an input array d_in and save the content in revers order into d_out. The revers is performed into the GPU.
 - d_in is [100, 110, 200, 220, 300]
then d_out must be [300, 220, 200, 110, 100]
 - **blockDim.x** is the number of threads per block
 - **gridDim.x** is the number of blocks in a grid
- Implement a Matrix Transpose using threads and blocks
- Implement a Matrix Multiplication for Matrix sizes 2048^2 .
Use max 512 threads x block.