

# PARALLEL PROGRAMMING II

# The MPI\_BARRIER

Blocks until all processes have reached this routine

```
INCLUDE 'mpif.h'
```

```
MPI_BARRIER(COMM, IERROR)
```

```
INTEGER    COMM, IERROR
```

# MPI DATA TYPES

<u>MPI datatype handle</u>	<u>C datatype</u>
MPI_INT	int
MPI_SHORT	short
MPI_LONG	long
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_CHAR	char
MPI_BYTE	unsigned char

# Calling MPI\_REDUCE

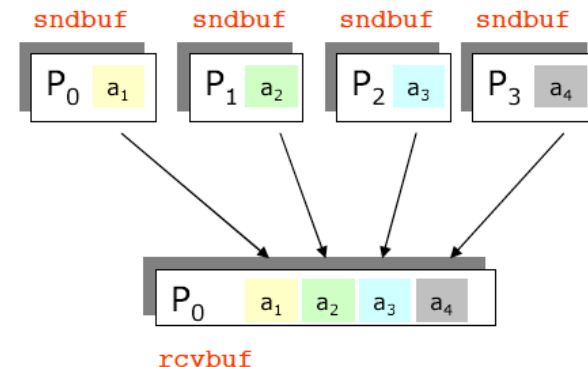
**MPI\_REDUCE(in, out, count, type, op, receiver, comm, err)**

in:	data to be sent (from all)
out:	storage for reduced data (on receiver)
count:	number of data items to be reduced
type:	type (=size) of data items
op:	reduction operation, e.g. <b>MPI_SUM</b>
receiver:	rank of sending processor of data
communicator:	group identifier, <b>MPI_COMM_WORLD</b>
err:	error status or <b>MPI_SUCCESS</b>

# MPI\_Gather

One-to-all communication: different data collected by the root process, from all others processes in the communicator.

It is the opposite of Scatter



**MPI\_GATHER( sendbuf, sendcount, sendtype, rcvbuf, rcvcount, recvttype, root, comm)**

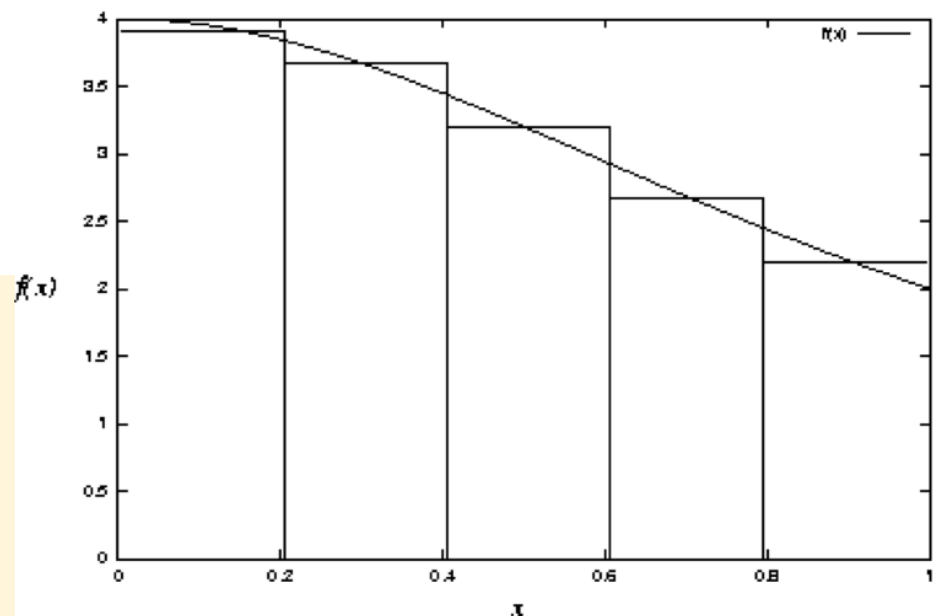
- [ IN sendbuf] starting address of send buffer (choice)
- [ IN sendcount] number of elements in send buffer (integer)
- [ IN sendtype] data type of send buffer elements (handle)
- [ OUT rcvbuf] address of receive buffer (choice, significant only at **root**)
- [ IN rcvcount] number of elements for any single receive (integer, significant only at **root**)
- [ IN recvttype] data type of recv buffer elements (significant only at **root**) (handle)
- [ IN root] rank of receiving process (integer)
- [ IN comm] communicator (handle)

# Approximate PI Using MPI collectives

$$\int_0^1 \frac{1}{1+x^2} dx = \arctan(x) \Big|_0^1 = \arctan(1) - \arctan(0) = \arctan(1) = \frac{\pi}{4}$$

$$\pi = 4 \int_0^1 \frac{1}{1+x^2} dx$$

Integrate, i.e determine area under function numerically using slices of  $h * f(x)$  at midpoints



```
#include <stdio.h>

int main(){
    long n , i ;
    double w,x,sum,pi,f,a;

    n = 100000000;
    w = 1.0/n;
    sum = 0.0;

    for ( i = 1 ; i <= n ; i++ ) {
        x = w * (i - 0.5);
        sum = sum + (4.0 / (1.0 + x * x ) );
    }

    pi = w * sum ;
    printf("Value of pi: %.16g\n", pi);

    return 0;
}
```

# Assignment

- 1) Implement the PI approximation in parallel using the Message Passing paradigm



# STANDARD BLOCKING SEND - RECV

**MPI\_SEND(buf, count, type, dest, tag, comm, ierr)**

**MPI\_RECV(buf, count, type, dest, tag, comm, status, ierr)**

**Buf** array of MPI type **type**.

**Count** (INTEGER) number of element of **buf** to be sent/recv

**Type** (INTEGER) MPI type of **buf**

**Dest** (INTEGER) rank of the destination process

**Tag** (INTEGER) number identifying the message

**Comm** (INTEGER) communicator of the sender and receiver

\* **Status** (INTEGER) array of size **MPI\_STATUS\_SIZE** containing communication status information

**ierr** (INTEGER) error code

*\* used only for receive operations*

# Wildcards

Both in Fortran and C **MPI\_RECV** accept wildcard:

- To receive from any source: **MPI\_ANY\_SOURCE**
- To receive with any tag: **MPI\_ANY\_TAG**
- Actual source and tag are returned in the receiver's status parameter => **status.MPI\_SOURCE**, **status.MPI\_TAG**

## **MPI\_GET\_COUNT(status, datatype, count)**

[ IN status] return status of receive operation (Status)

[ IN datatype] datatype of each receive buffer entry (handle)

[ OUT count] number of received entries (integer)

PROGRAM send\_recv

INCLUDE 'mpif.h'

INTEGER :: ierr, myid, nproc, **status(MPI\_STATUS\_SIZE)**

REAL A(2)

CALL MPI\_INIT(ierr)

CALL MPI\_COMM\_SIZE(MPI\_COMM\_WORLD, nproc, ierr)

CALL MPI\_COMM\_RANK(MPI\_COMM\_WORLD, myid, ierr)

**IF( myid .EQ. 0 ) THEN**

A(1) = 3.0

A(2) = 5.0

CALL MPI\_SEND(A, 2, MPI\_REAL, 1, 10, MPI\_COMM\_WORLD, ierr)

**ELSE IF( myid .EQ. 1 ) THEN**

CALL MPI\_RECV(A, 2, MPI\_REAL, 0, 10, MPI\_COMM\_WORLD, status, ierr)

WRITE(6,\*) myid,': a(1)=',a(1),' a(2)=',a(2)

**END IF**

CALL MPI\_FINALIZE(ierr)

END

## NON-BLOCKING SEND - RECV

**MPI\_ISEND(buf, count, type, dest, tag, comm, request, ierr)**

**MPI\_Irecv(buf, count, type, dest, tag, comm, request, ierr)**

**Buf** array of MPI type **type**.

**Count** (INTEGER) number of element of **buf** to be sent/recv

**Type** (INTEGER) MPI type of **buf**

**Dest** (INTEGER) rank of the destination process

**Tag** (INTEGER) number identifying the message

**Comm** (INTEGER) communicator of the sender and receiver

**Request** (INTEGER) request handler, used for checking the communication status

**lerr** (INTEGER) error code

# No-Blocking Checkpoint

## **MPI\_WAIT(request, status, ierr)**

**Request** (INTEGER) request handler, used for checking the communication status

**Status** (INTEGER) array of size **MPI\_STATUS\_SIZE** containing communication status information

**ierr** (INTEGER) error code

Wait until the communication handled by the object request is terminated. For test only use MPI\_TEST, for checkpoint of many communication use MPI\_WAITALL



# Assignments

- 1) Implement data exchange between two processes.  
Perform a test using 100 elements and 100000000 elements.

# Static Data Partitioning

**The simplest data decomposition schemes for dense matrices are 1-D block distribution schemes.**

row-wise distribution

$P_0$
$P_1$
$P_2$
$P_3$
$P_4$
$P_5$
$P_6$
$P_7$

column-wise distribution

$P_0$	$P_1$	$P_2$	$P_3$	$P_4$	$P_5$	$P_6$	$P_7$
-------	-------	-------	-------	-------	-------	-------	-------

# Distributed Data Vs Replicated Data

- Replicated data distribution is useful if it helps to reduce the communication among process at the cost of bounding scalability
- Distributed data is the ideal data distribution but not always applicable for all data-sets
- Usually complex application are a mix of those techniques





# Global Vs Local Indexes

In sequential code you always refer to global indexes

With distributed data you must handle the distinction between global and local indexes (and possibly implementing utilities for transparent conversion)

Local Idx

1	2	3
---	---	---

1	2	3
---	---	---

1	2	3
---	---	---

---

Global Idx

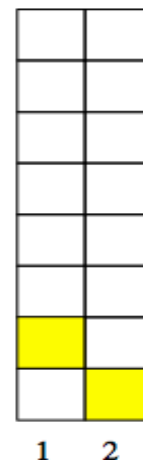
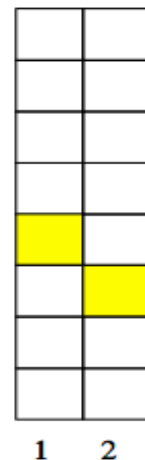
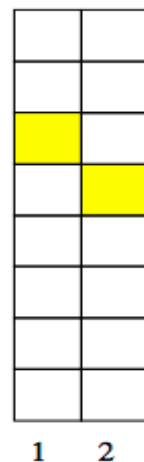
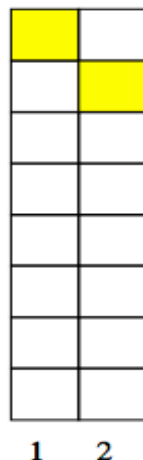
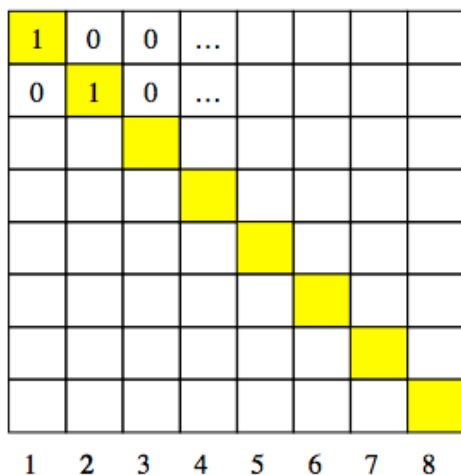
1	2	3
---	---	---

4	5	6
---	---	---

7	8	9
---	---	---



# Collaterals to Domain Decomposition /1

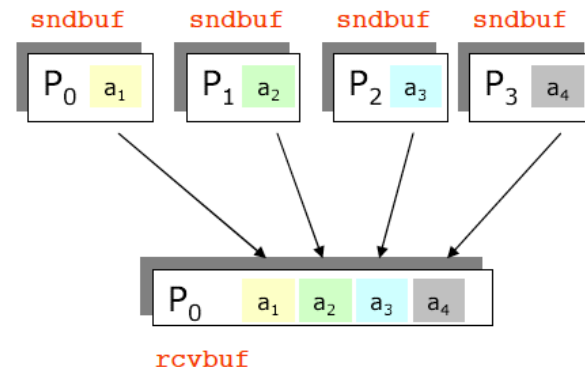


**Are all the domain's dimensions  
always multiple of the number  
of tasks/processes we are  
willing to use?**

# MPI\_Gatherv

One-to-all communication: different data collected by the root process, from all others processes in the communicator.

Messages can have different sizes and displacements



**MPI\_GATHERV( sndbuf, sendcount, sendtype, rcvbuf, recvcounts, displs, recvtype, root, comm)**

- [ IN sndbuf] starting address of send buffer (choice)
- [ IN sendcount] number of elements in send buffer (integer)
- [ IN sendtype] data type of send buffer elements (handle)
- [ OUT rcvbuf] address of receive buffer (choice, significant only at **root**)
- [ IN recvcounts] integer array (of length group size) containing the number of elements that are received from each process (significant only at **root**)
- [ IN displs] integer array (of length group size). Entry *i* specifies the displacement relative to rcvbuf at which to place the incoming data from process *i* (significant only at **root**)
- [ IN recvtype] data type of recv buffer elements (significant only at **root**) (handle)
- [ IN root] rank of receiving process (integer)
- [ IN comm] communicator (handle)

# Assignments

- 1) Initialize the Identity Matrix of size  $N \times N$  on distributed data. Consider to implement the option of printing the matrix on standard output, if a DEBUG mode is activated and the Matrix size  $\leq 10$ . Otherwise print the Matrix on a binary file.
- 2) Work initially on replicated data. Then optimize the solution for distributed data

# External MPI Resources

Here are some links to tutorials and literature:

CI-Tutor at NCSA: <http://www.citutor.org/>

MPI reference and mini tutorial at LLNL:

<http://computing.llnl.gov/tutorials/mpi/>

Designing and Building // Programs, by Ian Foster:

<http://www.mcs.anl.gov/~itf/dbpp/>

MPI standards: <http://www.mpi-forum.org/>

OpenMPI: <http://www.open-mpi.org>

MPICH: <http://www.mcs.anl.gov/research/projects/mpich2>