

# Shared Memory Programming Paradigm

**Ivan Girotto – [igirotto@ictp.it](mailto:igirotto@ictp.it)**


Information & Communication Technology Section (ICTS)  
International Centre for Theoretical Physics (ICTP)



Scuola Internazionale Superiore  
di Studi Avanzati





# Multi-CPUs & Multi-cores NUMA system



**Main Memory**

**Dual Socket (Westmere) - 24GB RAM**

# Processes and Threads



# TEXT

compiled code (a.out)

## DATA

## SHARED MEMORY

## STACK

System
env argv argc
auto variables for main()
auto variables for func()
<i>available for stack growth</i>
malloc.o (lib*.so)
printf.o (lib*.so)
<i>available for heap growth</i>
Heap (malloc arena)
global variables
"...%d..."
malloc.o (lib*.a)
printf.o (lib*.a)
file.o
main.o      func(72,73)
crt0.o (startup routine)

## High memory

← mfp – frame pointer (for main)

← stack pointer  
(grows downward if func() calls another function)

↑ library functions if  
dynamically linked  
(usual case)

← brk point

uninitialized data (bss)

initialized data

↑ library functions if  
statically linked  
(not usual case)

← ra (return address)

## Low memory

Stack illustrated after the call  
func(72,73) called from main(),  
assuming func defined by:  
func(int x, int y) {  
    int a;  
    int b[3];  
    /\* no other auto variables \*/  
}

Assumes int = long = char \* of  
size 4 and assumes stack at high  
address and descending down.

## Expanded view of the stack

Stack		Contents
Offset from current frame pointer (for func())	main() auto variables	
+12	73	y
+8	72	x
+4	ra	return address
0	mfp	caller's frame pointer
-4	garbage	a
-8	garbage	b[2]
-12	garbage	b[1]
-16	garbage	b[0]

All auto variables and parameters  
are referenced via offsets from the  
frame pointer.

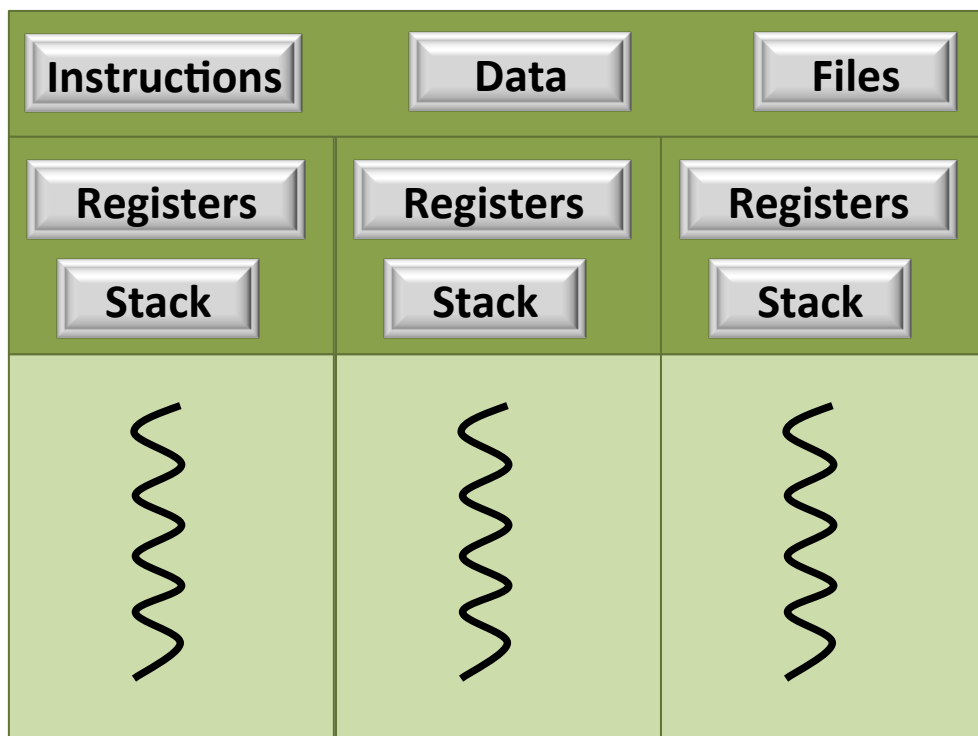
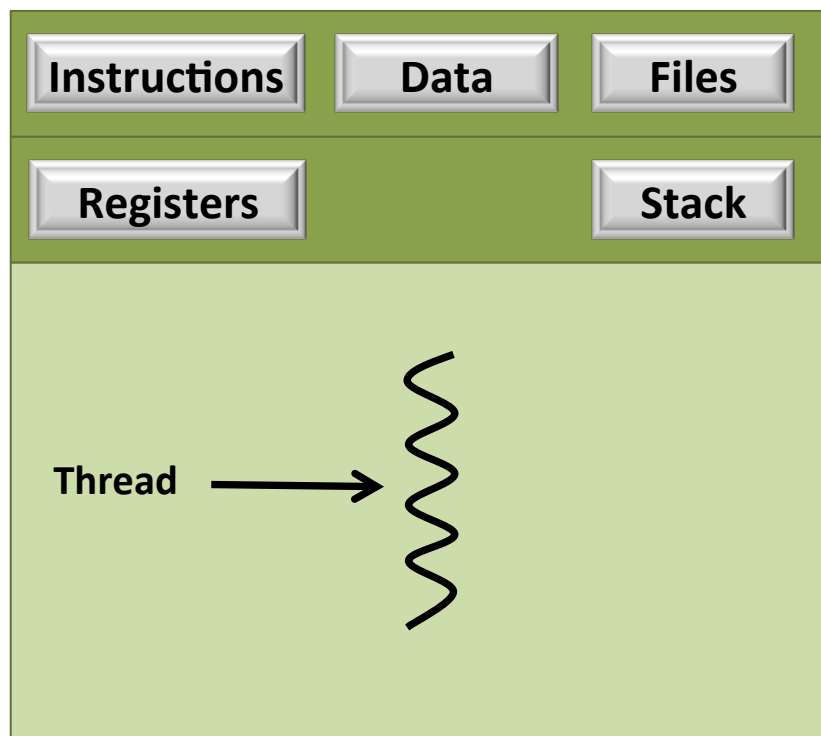
The frame pointer and stack pointer  
are in registers (for fast access).

When func returns, the return value  
is stored in a register. The stack pointer  
is move to the y location, the code  
is jumped to the return address (ra),  
and the frame pointer is set to mfp  
(the stored value of the caller's frame  
pointer). The caller moves the return  
value to the right place.

**EBP**

**ESP**

# Processes and Threads



# Multi-threading - Recap

- A thread is a (**lightweight**) process - an instance of a program plus its own data (private memory)
- Each thread can follow its own flow of control through a program
- Threads can share data with other threads, but also have private data
- Threads communicate with each other via the shared data.
- A *master thread* is responsible for co-ordinating the threads group

# OpenMP<sup>TM</sup>

# OpenMP (*Open spec. for Multi Processing*)

OpenMP **is not a computer language**

- Rather it works in conjunction with existing languages such as standard Fortran or C/C++

Application Programming Interface (API)

- that provides a **portable** model for parallel applications
- Three main components:
  - Compiler **directives**
  - **Runtime library** routines
  - **Environment variables**





# OpenMP Parallelization

OpenMP is directive based

- code (can) work without them

OpenMP can be added incrementally

OpenMP only works in shared memory

- multi-socket nodes, multi-core processors


OpenMP hides the calls to a threads library

- less flexible, but much less programming

**Caution:** write access to shared data can easily lead to race conditions and incorrect data



# OpenMP Parallelization



- Thread-based Parallelism
- Explicit Parallelism
- Fork-Join Model
- Compiler Directive Based
- Dynamic Threads

# Getting Started with OpenMP

OpenMP's constructs fall into 5 categories:

- Parallel Regions
- Work sharing
- Data Environment (scope)
- Synchronization
- Runtime functions/environment variables

OpenMP is essentially the same for both Fortran and C/C++



# Directives Format

A directive is a special line of source code with meaning only to certain compilers.

A directive is distinguished by a sentinel at the start of the line.

OpenMP sentinels are:

- Fortran: **!\$OMP** (or **C\$OMP** or **\*\$OMP**)
- C/C++: **#pragma omp**

# OpenMP: Parallel Regions

For example, to create a 4-thread parallel region:  
each thread calls `foo(ID,A)` for **ID** = **0** to **3**

Each thread redundantly  
executes the code within  
the structured block

thread-safe routine: A routine that performs  
the intended function even when executed  
concurrently (by more than one thread)

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID =omp_get_thread_num();  
    foo(ID,A);  
}  
printf( "All Done\n" );
```

double A[1000];

omp\_set\_num\_threads(4);

A single copy of A is shared between all threads.

foo(0,A);    foo(1,A);    foo(2,A);    foo(3,A);

printf( "All Done\n"

Threads wait here for all threads to finish before proceeding (i.e. barrier).

## How many threads?

- The number of threads in a parallel region is determined by the following factors:
  - Use of the `omp_set_num_threads()` library function
  - Setting of the `OMP_NUM_THREADS` environment variable
  - The implementation **default**
  - Threads are numbered from 0 (master thread) to N-1.

# OpenMP runtime library

`OMP_GET_NUM_THREADS()` – returns the current # of threads.

`OMP_GET_THREAD_NUM()` - returns the id of this thread.

`OMP_SET_NUM_THREADS(n)` – set the desired # of threads.

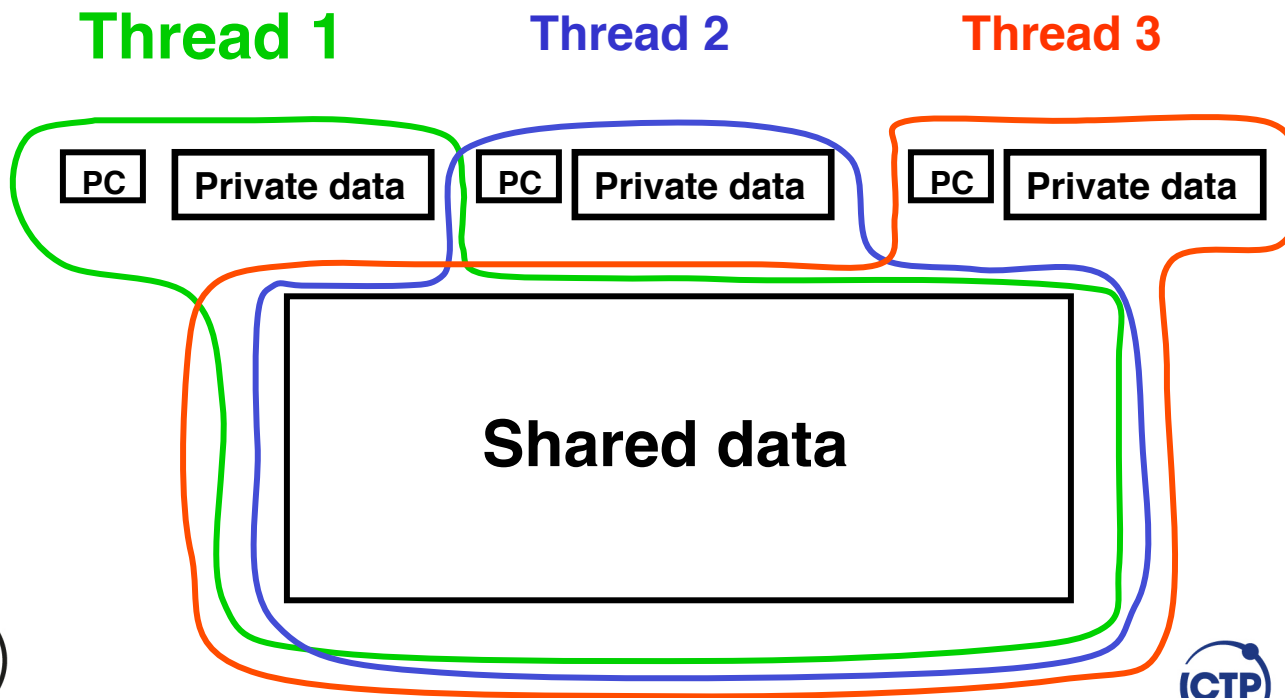
`OMP_IN_PARALLEL()` – returns .true. if inside parallel region.

`OMP_GET_MAX_THREADS()` - returns the # of possible threads.





# Memory footprint



Program



Thread 1

```
load a
add a 1
store a
```


Thread 2

```
load a
add a 1
store a
```

Private  
data



Shared  
data



# Simple C OpenMP Program

```
#include <omp.h>
#include <stdio.h>

int main ( ) {

    printf("Starting off in the sequential world.\n");
    #pragma omp parallel
    {
        printf("Hello from thread number %d\n", omp_get_thread_num() );
    }
    printf("Back to the sequential world.\n");
    return 0;
}
```

# Variable Scooping

All existing variable still exist inside a parallel region

- by default SHARED between all threads

But work sharing requires private variables

- PRIVATE clause to OMP PARALLEL directive
- Index variable of a worksharing loop
- All declared local variable within a parallel region
- The FIRSTPRIVATE clause would initialize the private instances with the contents of the shared instance

Be aware of the sharing nature of static variables



# Exploiting Loop Level Parallelism

Loop level Parallelism: parallelize only loops

Easy to implement

Highly readable code

Less than optimal performance (sometimes)

Most often used

## Parallel Loop Directives

Fortran do loop directive

- `!$omp do`

C\C++ for loop directive

- `#pragma omp for`

These directives do not create a team of threads but assume there has already been a team forked.

If not inside a parallel region shortcuts can be used.

- `!$omp parallel do`
- `#pragma omp parallel for`

## Parallel Loop Directives /2

These are equivalent to a parallel construct followed immediately by a worksharing construct.

**!\$omp parallel do**

**Same as**

**!\$omp parallel**

**...**

**!\$omp do**

**#pragma omp parallel for**

**Same as**

**#pragma omp parallel**

**...**

**#pragma omp for**

```
integer :: N, start, len, numth, tid, i, end
double precision, dimension (N) :: a, b, c
!$OMP PARALLEL PRIVATE (start, end, len, numth, tid, i)
  numth = omp_get_num_threads()
  tid = omp_get_thread_num()
  len = N / numth
  if( tid .lt. mod( N, numth ) ) then
    len = len + 1
    start = len * tid + 1
  else
    start = len * tid + mod( N, numth ) + 1
  endif
  end = start + len - 1
  do i = start, end
    a(i) = b(i) + c(i)
  end do
!OMP END PARALLEL
```

**Not the intended  
mode for OpenMP**



# How is OpenMP Typically Used?

OpenMP is usually used to parallelize loops:

**Split-up this loop between multiple threads**


```
void main()
{
    double Res[1000];

    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Sequential program

```
void main()
{
    double Res[1000];
    #pragma omp parallel for
    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Parallel program



# Work-Sharing Constructs

Divides the execution of the enclosed code region among the members of the team that encounter it.

Work-sharing constructs **do not** launch new threads.

No implied barrier upon entry to a work sharing construct.

However, there is an implied barrier at the end of the work sharing construct (unless **nowait** is used).



# Work Sharing Constructs - example

Sequential code

```
for(i=0;I<N;i++) { a[i] = a[i] + b[i];}
```


```
#pragma omp parallel
{
    int id, i, Nthrds, istart, iend;
    id = omp_get_thread_num();
    Nthrds = omp_get_num_threads();
    istart = id * N / Nthrds;
    iend = (id+1) * N / Nthrds;
    for(i=istart;I<iend;i++) {a[i]=a[i]+b[i];}
}
```

OpenMP // Region

```
#pragma omp parallel
#pragma omp for schedule(static)
for(i=0;I<N;i++) { a[i]=a[i]+b[i];}
```

OpenMP Parallel  
Region and a work-  
sharing for construct

# schedule(static [,chunk])




- Iterations are divided evenly among threads
- If chunk is specified, divides the work into chunk sized parcels
- If there are N threads, each thread does every N<sup>th</sup> chunk of work.

```
!$OMP PARALLEL DO &
!$OMP SCHEDULE(STATIC,3)
```

```
DO J = 1, 36
    Work (j)
END DO
```

```
!$OMP END DO
```

# schedule(dynamic [,chunk])



- Divides the workload into chunk sized parcels.
- As a thread finishes one chunk, it grabs the next available chunk.
- Default value for chunk is one.
- More overhead, but potentially better load balancing.

```
!$OMP PARALLEL DO & !
$OMPSCHEDULE(DYNAMIC,1)
```

```
DO J = 1, 36
    Work (j)
END DO
```

```
!$OMP END DO
```

# The Schedule Clause SCHEDULE (type [,chunk])

The schedule clause effects how loop iterations are mapped onto threads

`schedule(static [,chunk])`

- Deal-out blocks of iterations of size “chunk” to each thread

`schedule(dynamic [,chunk])`

- Each thread grabs “chunk” iterations off a queue until all iterations have been handled

`schedule(guided [,chunk])`

- Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds

`schedule(runtime)`

- Schedule and chunk size taken from the OMP\_SCHEDULE environment variable

## No Wait Clauses

- No wait: if specified then threads do not synchronise at the end of the parallel loop.

For Fortran, the END DO directive is optional with NO WAIT being the default.

Note that the nowait clause is incompatible with a simple parallel region meaning that using the composite directives will not allow you to use the nowait clause.



## OpenMP: Reduction(op : list)

The variables in “list” must be shared in the enclosing parallel region.

Inside a parallel or a worksharing construct:

- A local copy of each list variable is made and initialized depending on the “op” (e.g. 0 for “+”)
- pair wise “op” is updated on the local value
- Local copies are reduced into a single global copy at the end of the construct.

# OpenMP: A Reduction Example

```
#include <omp.h>
#define NUM_THREADS 2
void main ()
{
    int i;
    double ZZ, func(), sum=0.0;


    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel for reduction(+:sum) private(ZZ)
    for (i=0; i< 1000; i++){
        ZZ = func(i);
        sum = sum + ZZ;
    }
}
```

# Compute PI

$$\int_0^1 \frac{1}{1+x^2} dx = \arctan(x) \Big|_0^1 = \arctan(1) - \arctan(0) = \arctan(1) = \frac{\pi}{4}$$

$$\pi = 4 \int_0^1 \frac{1}{1+x^2} dx$$

Integrate, i.e determine area under function numerically using slices of  $h * f(x)$  at midpoints



```
program calc_pi
  integer :: i
  real(kind=8),parameter ::PIREF = 3.141592653589793d0
  real(kind=8) :: w,x,s,pi
  integer,parameter :: n = 1000000000
  w=1.0/n
  s=0.0
  !$OMP parallel do private(i,x) reduction(+:s)
  do i=1,n
    x = w * (i - 0.5)
    s = s + 4.0 / (1.0 + x*x)
  end do
  pi = w * s
  print *, 'pi is ', pi, 'Error is ', abs(pi - PIREF)
end program calc_pi
```

## if CLAUSE

We can make the parallel region directive itself conditional.

Fortran: **IF** (*scalar logical expression*)

C/C++: **if** (*scalar expression*)

```
#pragma omp parallel if (tasks > 1000)
{
    while(tasks > 0) donexttask();
}
```

# SYNCHRONIZATION

# OpenMP: How do Threads Interact?

OpenMP is a shared memory model.

- Threads communicate by **sharing variables**.

Unintended sharing of data can lead to **race conditions**:

- race condition: when the program's outcome changes as the threads are scheduled differently.

To control race conditions:

- Use **synchronization** to protect data conflicts.

Synchronization is expensive so:

- Change how data is stored to minimize the need for synchronization.

Note that updates to shared variables:

(e.g.  $a = a + 1$ )

are *not* **atomic**!

If two threads try to do this at the same time, one of the updates may get overwritten.



Program



Thread 1

```
load a
add a 1
store a
```


Thread 2

```
load a
add a 1
store a
```

Private  
data



Shared  
data



## Barrier

- |         |                              |
|---------|------------------------------|
| Fortran | - <b>!\$OMP BARRIER</b>      |
| C\C++   | - <b>#pragma omp barrier</b> |

This directive synchronises the threads in a team by causing them to wait until all of the other threads have reached this point in the code.

Implicit barriers exist after work sharing constructs. The `nowait` clause can be used to prevent this behaviour.



Add a note about single/master

## Critical

Only one thread at a time can enter a **critical** section.

Example: pushing and popping a task stack

```
!$OMP PARALLEL SHARED(STACK),PRIVATE(INEXT,INEW)
```

...

```
!$OMP CRITICAL (STACKPROT)
```

```
  inext = getnext(stack)
```

```
!$OMP END CRITICAL (STACKPROT)
```

```
  call work(inext,inew)
```

```
!$OMP CRITICAL (STACKPROT)
```

```
  if (inew .gt. 0) call putnew(inew,stack)
```

```
!$OMP END CRITICAL (STACKPROT)
```

...

```
!$OMP END PARALLEL
```



## Atomic

**Atomic** is a special case of a critical section that can be used for certain simple statements

Fortran: **!\$OMP ATOMIC**  
*statement*

where *statement* must have one of these forms:

$x = x \text{ op } \text{expr}, \quad x = \text{expr op } x, \quad x = \text{intr} (x, \text{expr}) \text{ or}$   
 $x = \text{intr}(\text{expr}, x)$

*op* is one of +, \*, -, /, .and., .or., .eqv., Or .neqv.

*intr* is one of MAX, MIN, IAND, IOR or IEOR




# Non Parallelizzabile

Show an example of Instruction dependency

# Poor Performances

Often naive approaches multi-threaded programming results in poor performances

- Modern NUMA architecture requires specific attention, specially considering multithreaded-programming
- The overhead of thread-management is not always negligible
- Reduce to minimum the critical regions
- FALSE SHARING is behind the corner
- Anything shared is a possible source of race condition (if write access)



# Exercises

1. write an “hello-world” program that prints on std output how many threads are executed and the thread\_ID for each thread
2. parallelize the heat equation code using OpenMP
3. parallelize the fast-transpose using OpenMP
4. parallelize the code provided in class using OpenMP

Note: perform performance analysis for the points 2-4. Write a Makefile that somehow allows to compile the serial and the OpenMP versions of the code. Instrument the code to print at runtime the number of threads, in case of a parallel version.